Getting Started with
Machine Learning using Spark in Azure HDInsight

Ingyu Lee

Troy University

https://github.com/Azure-Bootcamp-Troy/SparkOnHDInsight

Content

- What is **Spark** on **HDInsight**?
- Demo: Create Spark Cluster on Azure

- What is an **RDD, DataFrames, Spark SQL**?

- Demo: Operations of RDD, DataFrames and SQL

- What is **MLlib**?

- Demo: Machine Learning using Spark

- For Whom this talk
  - Business users to analyze Big Data using Spark on Azure Cluster.

- Prerequisite
  - None

References

This talk is based on this course

Many excellent books released in the past year or two!

Advanced Analytics with Spark (2015), written by
Sandy Ryza, Uri Laserson, Sean Owen, and Josh Wills

Big Data Analysis with Apache Spark
Learn how to apply data science techniques using parallel programming in Apache Spark to explore big data.

Distributed Machine Learning with Apache Spark
Learn the underlying principles required to develop scalable machine learning pipelines and gain hands-on experience using Apache Spark.

Introduction to Apache Spark
Learn the fundamentals and architecture of Apache Spark, the leading cluster-computing framework among professionals.
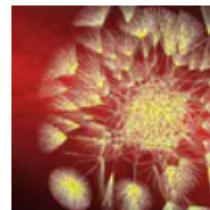
Microsoft
Implementing Predictive Analytics with Spark in

Microsoft
Processing Real-Time Data Streams in Azure

Microsoft
Processing Big Data with Hadoop in Azure

Big Data for Data Engineers
Yandex

Learn More

Big Data
University of California, San Diego

Learn More

HDInsight – What is it?

A standard Apache Hadoop distribution offered as a managed service on Microsoft Azure

- Based on Hortonworks Data Platform (HDP)
- Provisioned as clusters on Azure that can run on Windows or Linux servers
- Offers capacity-on-demand, pay-as-you-go pricing model
- Integrates with:
  - Azure Blob Storage and Azure Data Lake Store for Hadoop File System (HDFS)
  - Azure Portal for management and administration
  - Visual Studio for application development tooling

In addition to the core, HDInsight supports the Hadoop ecosystem

What is Apache Spark?

General, open-source cluster computing engine

Well-suited for machine learning
- Fast iterative procedures
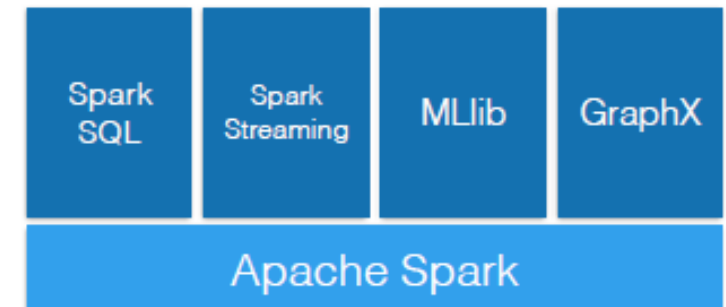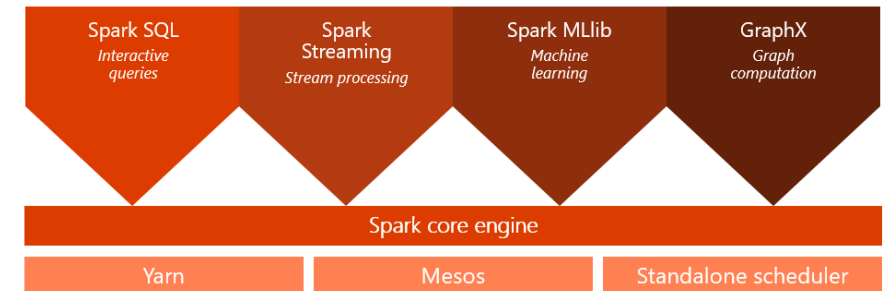- Efficient communication primitives

Simple and Expressive
- APIs in Scala, Java, Python, R
- Interactive Shell

Integrated Higher-Level Libraries

Apache Spark: A unified framework

A unified, open source, parallel data processing framework for big data analytics

Spark SQL *Interactive queries*	Spark Streaming *Stream processing*	Spark MLlib *Machine learning*	GraphX *Graph computation*

Spark core engine

Yarn	Mesos	Standalone scheduler

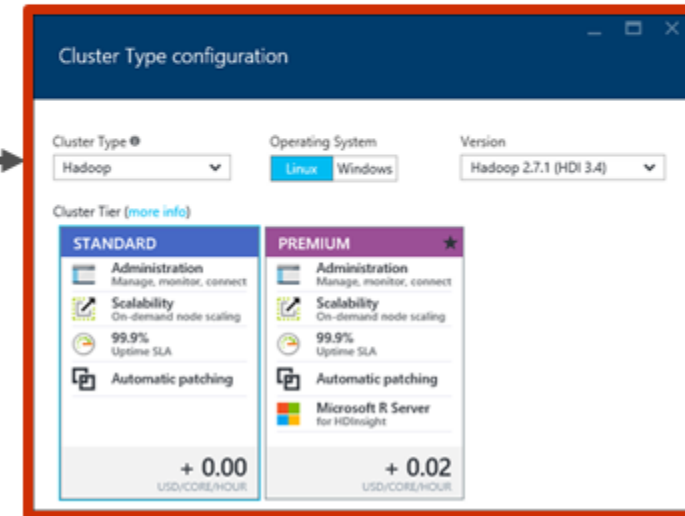| Spark SQL	Spark Streaming	MLlib	GraphX

Apache Spark

Demo: Creating an HDInsight Spark cluster
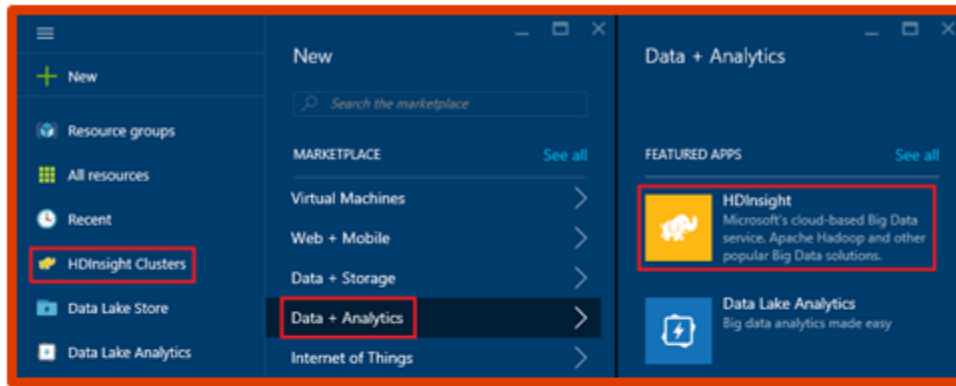
Apache Spark use cases

The Azure console lists all types of HDInsight clusters (HBase, Storm, and Spark) currently provisioned.

Demo: Ambari Dashboard

The Ambari Dashboard provides links to access:

Ambari web UI – monitors Spark clusters
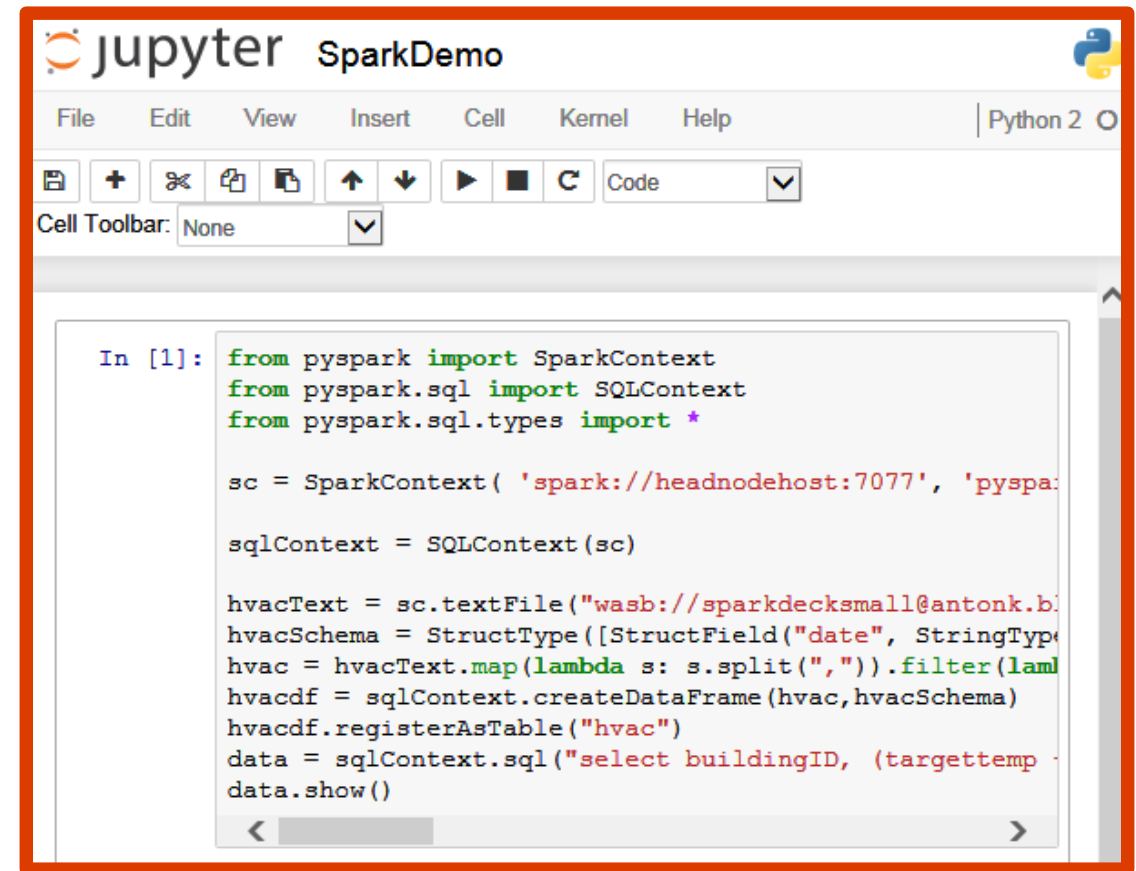
Hosts (Resource Manager) – controls amount of resources allocated to various Spark cluster components

Notebooks – interactive web-based tools to develop and run Spark programs
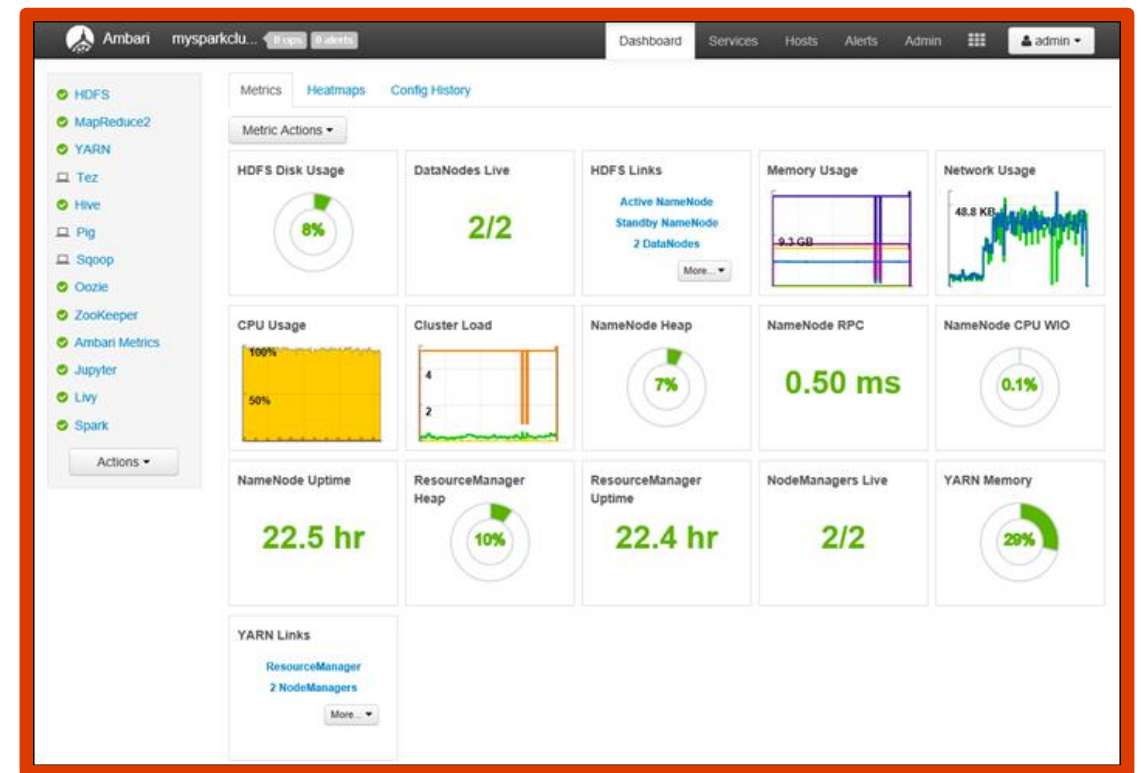
Demo: Jupyter notebooks

- Web-based interactive consoles for
  - Experimentation
  - Collaboration

- Spark HDInsight clusters include Jupyter
  - Interactive Python
  - Interactive Scala

What is RDD?

- The core abstraction for data in Spark is the *resilient distributed dataset* (RDD)

- An RDD represents a collection of items that can be distributed across compute nodes

- APIs for working with RDDs are provided for Java, Python, and Scala
  - HDInsight distribution includes Python and Scala shells

Dataset	Distributed	Resilient
Data storage created from: HDFS, S3, HBase, JSON, text, Local hierarchy of folders	Distributed across the cluster of machines	Recover from errors, e.g. node failure, slow processes
Or created transforming another RDD	Divided in partitions, atomic chunks of data	Track history of each partition, re-run

Resilient Distributed Datasets or RDDs address this by enabling fault-tolerant, distributed, in-memory computations.

RDD Operations

RDDs: Transformations and actions

Some Transformations

Transformation	Description
map(*func*)	return a new distributed dataset formed by passing each element of the source through a function *func*
filter(*func*)	return a new dataset formed by selecting those elements of the source on which *func* returns true
distinct([*numTasks*]))	return a new dataset that contains the distinct elements of the source dataset
flatMap(*func*)	similar to map, but each input item can be mapped to 0 or more output items (so *func* should return a Seq rather than a single item)

Key-Value Transformation	Description
reduceByKey(*func*)	return a new distributed dataset of (K,V) pairs where the values for each key are aggregated using the given reduce function *func*, which must be of type (V,V) → V
sortByKey()	return a new dataset (K,V) pairs sorted by keys in ascending order
groupByKey()	return a new dataset of (K, Iterable<V>) pairs

Transformations	RDD	Actions	Value

Some Actions

Action	Description
reduce(*func*)	aggregate dataset's elements using function *func*. *func* takes two arguments and returns one, and is commutative and associative so that it can be computed correctly in parallel
take(*n*)	return an array with the first *n* elements
collect()	return all the elements as an array WARNING: make sure will fit in driver program
takeOrdered(*n*, *key=func*)	return n elements ordered in ascending order or as specified by the optional key function

Demo: RDD

RDDs can be created from stable storage or by transforming other RDDs.

```
In [ ]:  fruits = spark.sparkContext.textFile('wasb:///example/data/fruits.txt')
         yellowThings = spark.sparkContext.textFile('wasb:///example/data/yellowthings.txt')
```

Transformations create a new dataset from an existing one. Transformations are lazy, meaning that no transformation is executed until you execute an action.

```
# map
fruitsReversed = fruits.map(lambda fruit: fruit[::-1])
```

```
# filter
shortFruits = fruits.filter(lambda fruit: len(fruit) <= 5)
```

```
# flatMap
characters = fruits.flatMap(lambda fruit: list(fruit))
```

```
# union
fruitsAndYellowThings = fruits.union(yellowThings)
```

```
# intersection
yellowFruits = fruits.intersection(yellowThings)
```

```
# distinct
distinctFruitsAndYellowThings = fruitsAndYellowThings.distinct()
distinctFruitsAndYellowThings
```

```
# groupByKey
yellowThingsByFirstLetter = yellowThings.map(lambda thing: (thing[0], thing)).groupByKey()
```

```
# reduceByKey
numFruitsByLength = fruits.map(lambda fruit: (len(fruit), 1)).reduceByKey(lambda x, y: x + y)
```

Actions return a value to the driver program after running a computation on the dataset.

```
# collect
fruitsArray = fruits.collect()
yellowThingsArray = yellowThings.collect()
fruitsArray
```

```
# count
numFruits = fruits.count()
numFruits
```

```
# take
first3Fruits = fruits.take(3)
first3Fruits
```

```
# reduce
letterSet = fruits.map(lambda fruit: set(fruit)).reduce(lambda x, y: x.union(y))
letterSet
```

What is DataFrames?

- A distributed collection of data organized into named columns.
- Similar to RDDs with schema.
- Conceptually equivalent to tables in relational database, or to DataFrames in R/Python.
- With domain-specific functions designed for common tasks:
  - Metadata
  - Sampling
  - Project, filter, aggregation, and join
  - UDFs

RDDs are a collection of opaque objects (such as internal structures unknown to Spark).

User	User	User
User	User	User

DataFrames is a collection of objects with schema that are known to Spark SQL..

Name	Age	Sex
Name	Age	Sex
Name	Age	Sex
Name	Age	Sex

DataFrame Operations

DataFrames provide a domain-specific language for structured data manipulation in Scala, Java, and Python.

```
>>> val df = sqlContext.jsonFile("somejasonfile.json")

>>> df.show()                               // Show the contents of the DataFrame

>>> df.printSchema()                        // Print the schema in a tree format

>>> df.select("name").show()                // Select and show the name columns

>>> df.select(df("name"),df ("age") +1).show()   // Select all but increment the age by 1

>>> df.filter(df("age") > 21).show()        // Select people older than 21

>>> df.groupBy("age").count().show()        // Count people by age
```

What is Spark SQL?

1 Create an Azure storage account	2 HDInsight makes Apache Spark available as a service in cloud.	3 Run Spark SQL statements using notebooks
HDInsight uses Azure Blob storage account for storing data.	HDInsight makes Apache Spark available as a service in cloud.	You run interactive Spark SQL statements using notebooks.

Tables and Queries

A DataFrame can be registered as a table that can then be used in SQL queries.

```
// First create a DataFrame from JSON file.
>>> val df = sqlContext.jsonFile(”Users.json“)

// Register the DataFrame as a temporary table. Temp tables exist only during lifetime of this SQLContext instance.
 >>> val usertable = sqlContext.registerDataFrameAsTable(df, “UserTable”)

// Alternatively, execute a SQL query on the table. The query returns a DataFrame.
>>> val teenagers = sqlContext.sql(“select Age as Years from UserTable where age > 13 and age <= 19”)
```
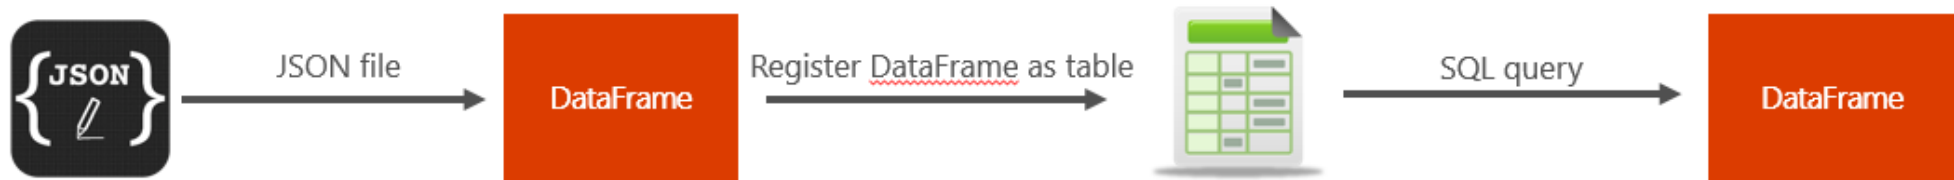
Spark SQL Operations

Spark RDD API

```
rdd
rdd1 = rdd.map(lambda x: x.split("\t"))
rdd2 = rdd1.map(lambda x: (x[0], x[2]))
```

VS

SQL on Spark

```
select user_id, url from access_log
```

- No data parsing
- Code optimization
- Syntax is easier
- No overhead

Demo: DataFrame and SQL

create a dataframe from a CSV file as shown below.

```
df = spark.read.csv('wasb:///HdiSamples/HdiSamples/SensorSampleData/building/building.csv', header=True, inferSchema=True)
```

```
# show the content of the dataframe
df.show()
```

```
# Print the dataframe schema in a tree format
df.printSchema()
```

```
# Create an RDD from the dataframe
dfrdd = df.rdd
dfrdd.take(3)
```

```
dfrdd = df.rdd
dfrdd.take(3)
```

```
# Retrieve a given number of rows from the dataframe
df.limit(3).show()
```

```
df.limit(3).show()
```

```
# Retrieve specific columns from the dataframe
df.select('BuildingID', 'Country').limit(3).show()
```

```
# Use GroupBy clause with dataframe
df.groupBy('HVACProduct').count().select('HVACProduct', 'count').show()
```

run SQL queries over dataframes once you register them as temporary tables within the SparkSession

```
# Register the dataframe as a temporary table called HVAC
df.registerTempTable('HVAC')
```

```
%%sql
SELECT * FROM HVAC WHERE BuildingAge >= 10
```

```
%%sql
SELECT BuildingID, Country FROM HVAC LIMIT 3
```

What is Machine Learning?

A Definition

Constructing and studying methods that learn from and make predictions on data

Broad area involving tools and ideas from various domains
- Computer Science
- Probability and Statistics
- Optimization
- Linear Algebra

Some Examples

Face recognition

Link prediction

Text or document classification, e.g., spam detection

Protein structure prediction

Games, e.g., Backgammon or Jeopardy

How does Machine Learning work?

- Data labels
- Direct feedback
- Predict outcome

Triangle diagram:
- Supervised learning (top)
- Machine Learning (center)
- Unsupervised learning (bottom left)
- Reinforcement learning (bottom right)

Supervised learning. Learning from labeled observations
- Labels 'teach' algorithm to learn mapping from observations to labels

Unsupervised learning. Learning from unlabeled observations
- Learning algorithm must find latent structure from features alone
- Can be goal in itself (discover hidden patterns, exploratory data analysis)
- Can be means to an end (preprocessing for supervised task)

Flow diagram:
Dataset → Features / Labels → Algorithm
Feature Vector → Trained Model → Expected Label
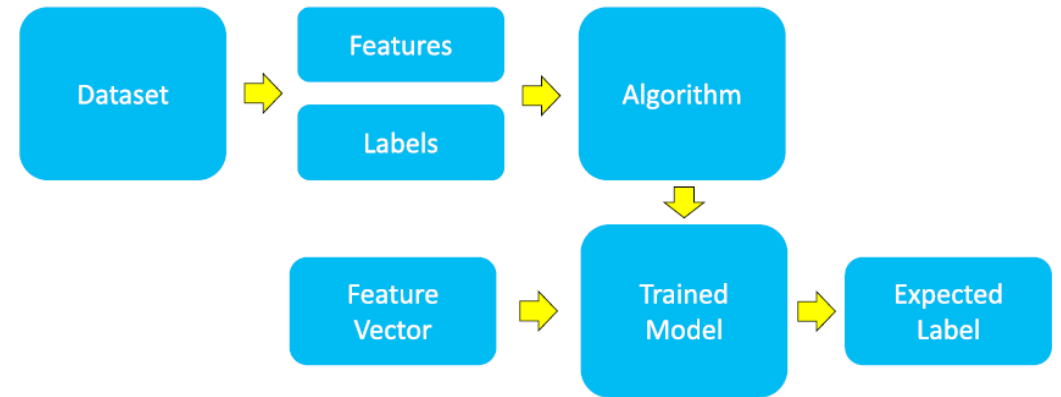
Observations. Items or entities used for learning or evaluation, e.g., emails

Features. Attributes (typically numeric) used to represent an observation, e.g., length, date, presence of keywords

Labels. Values / categories assigned to observations, e.g., spam, not-spam

Training and Test Data. Observations used to train and evaluate a learning algorithm, e.g., a set of emails along with their labels
- Training data is given to the algorithm for training
- Test data is withheld at train time

What is MLlib?

- MLlib is a collection of **machine learning algorithms optimized to run in a parallel, distributed manner on Spark clusters**.

- MLlib helps lead to better performance on large data sets.

- MLlib integrates seamlessly with other Spark components.

- MLlib applications are developed in Java, Scala, and Python.

Type	Algorithms
Supervised	Classification and regression: • Linear models (SVMs) logistic regression and linear regression • Naïve Bayes • Decision trees • Ensembles of trees (random forest, gradient-boosted trees) • Isotonic regression
Unsupervised	Clustering: • k-means and streaming k-means • Gaussian mixture • Power iteration clustering (PIC) • Latent Dirichlet Allocation (LDA)
Recommendation	Collaborative filtering: • Alternating least squares (ALS)

How to.. Machine Learning in Apache Spark?

- All primitives in Spark Machine Learning are *Vectors*

- *Features* are represented by a Vector

- Vectors can contain other Vectors and so be Dense or Sparse

- Spark uses *LabeledPoints* to encapsulate a Vector and a Label

- RDDs are transformed into Vectors through map functions

MLlib Terminology

ML: Transformer

- A *Transformer* is a class which can transform one DataFrame into another DataFrame
- A Transformer implements `transform()`

- Examples
  - HashingTF
  - LogisticRegressionModel
  - Binarizer

```
val tokenizer = new Tokenizer().setInputCol("text").setOutputCol("words")
val hashingTF = new
HashingTF().setNumFeatures(1000).setInputCol(tokenizer.getOutputCol).s
etOutputCol("features")
val lr = new LogisticRegression().setMaxIter(10).setRegParam(0.01)
val pipeline = new Pipeline().setStages(Array(tokenizer, hashingTF, lr))

val model = pipeline.fit(training)

val modelem = model.transform(test).select("id", "label", "text",
"probability", "prediction")
```

ML: Estimator

- An *Estimator* is a class which can take a DataFrame and produce a Transformer
- An Estimator implements `fit()`

- Examples
  - LogisticRegression
  - StandardScaler
  - Pipeline

```
val tokenizer = new Tokenizer().setInputCol("text").setOutputCol("words")
val hashingTF = new
HashingTF().setNumFeatures(1000).setInputCol(tokenizer.getOutputCol).s
etOutputCol("features")
val lr = new LogisticRegression().setMaxIter(10).setRegParam(0.01)
val pipeline = new Pipeline().setStages(Array(tokenizer, hashingTF, lr))

val model = pipeline.fit(training)

val modelem = model.transform(test).select("id", "label", "text",
"probability", "prediction")
```

What is Classification?

- To infer a function using labeled training data. The function can then be used **to determine the label for a new dataset.**

- Algorithms
  - Logistic regression
  - Decision Trees
  - Neural networks
  - Support vector machines
  - Naïve Bayes

Binary classification

Examples (telecommunication co.)
- **Churn**: Customers who have an inclination to switch to a different telecommunication provider.
- **Propensity to Buy**: Customers willing to buy new products or services.
- **Upselling**: Customers willing to buy upgraded services or add-ons.

Example: Classification I

Import Spark SQL and Spark ML Libraries

First, import the libraries you will need:

```python
from pyspark.sql.types import *
from pyspark.sql.functions import *

from pyspark.ml.classification import LogisticRegression
from pyspark.ml.feature import VectorAssembler
```

You will load this data into a DataFrame and display it.

Step 2: Load Source Data

```python
csv = spark.read.csv('wasb:///data/flights.csv', inferSchema=True, header=True)
csv.show()
```

```
+----------+---------+-------+---------------+-------------+--------+--------+
|DayofMonth|DayOfWeek|Carrier|OriginAirportID|DestAirportID|DepDelay|ArrDelay|
+----------+---------+-------+---------------+-------------+--------+--------+
|        19|        5|     DL|          11433|        13303|      -3|       1|
|        19|        5|     DL|          14869|
```

Step 3: Prepare the Data

```python
data = csv.select("DayofMonth", "DayOfWeek", "OriginAirportID", "DestAirportID", "DepDelay",
                  ((col("ArrDelay") > 15).cast("Int").alias("Late")))
data.show()
```

```
+----------+---------+---------------+-------------+--------+----+
|DayofMonth|DayOfWeek|OriginAirportID|DestAirportID|DepDelay|Late|
+----------+---------+---------------+-------------+--------+----+
|        19|        5|          11433|        13303|      -3|   0|
|        19|        5|          14869|        12478|       0|   0|
|        19|        5|          14057|        14869|      -4|   0|
```

Example: Classification II

Step 4: Split the Data

```python
splits = data.randomSplit([0.7, 0.3])
train = splits[0]
test = splits[1]
train_rows = train.count()
test_rows = test.count()
print "Training Rows:", train_rows, " Testing Rows:", test_rows
```

Step 5: Prepare the Training Data

```python
assembler = VectorAssembler(inputCols = ["DayofMonth", "DayOfWeek", "OriginAirportID", "DestAirportID", "DepDelay"],
                            outputCol="features")
training = assembler.transform(train).select(col("features"), col("Late").alias("label"))
training.show()
```

```
+--------------------+-----+
|            features|label|
+--------------------+-----+
|[1.0,1.0,10140.0,...|    0|
```

Step 6: Train a Classification Model

```python
lr = LogisticRegression(labelCol="label",featuresCol="features",maxIter=10,regParam=0.3)
model = lr.fit(training)
print "Model trained!"
```

```
Model trained!
```

Example: Classification III

Step 7: Prepare the Testing Data

```
testing = assembler.transform(test).select(col("features"), col("Late").alias("trueLabel"))
testing.show()
```

```
+--------------------+---------+
|            features|trueLabel|
+--------------------+---------+
|[1.0,1.0,10140.0,...|        0|
|[1.0,1.0,10140.0,...|        0|
```

Step 8: Test the Model

```
prediction = model.transform(testing)
predicted = prediction.select("features", "prediction", "probability", "trueLabel")
predicted.show(100)
```

```
+--------------------+----------+--------------------+---------+
|            features|prediction|         probability|trueLabel|
+--------------------+----------+--------------------+---------+
|[1.0,1.0,10140.0,...|       0.0|[0.83121028425635...|        0|
|[1.0,1.0,10140.0,...|       0.0|[0.82357817040192...|        0|
```

What is Regression?

- Regression techniques are used **to predict response variables with numerical outcomes**
  - E.g., predicting the miles per gallon of a car or predicting the temperature of a city.

- The input variables may be numeric or categorical. However, the output (or response variable) is numeric.

MPG vs. Horsepower

$$Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \beta_3 X_3 + \beta_4 X_4 + \cdots + \varepsilon$$

where

- Y is the response variable (the outcome you are trying to predict), such as miles per gallon.

- X_1, X_2, X_3, etc. are the independent variables used to predict the outcome.

- β_0 is a constant that is the intercept of the regression line.

- $\beta_1, \beta_2, \beta_3$, etc. are the coefficients of the independent variables. These refer to the partial slopes of each variable.

- ε is the error or noise associated with the response variable that cannot be explained by the independent variables X1, X2, and X3.

Example: Regression I

Import Spark SQL and Spark ML Libraries

First, import the libraries you will need:

```
from pyspark.sql.types import *
from pyspark.sql.functions import *

from pyspark.ml.regression import LinearRegression
from pyspark.ml.feature import VectorAssembler
```

Starting Spark application

Step 2: Load Source Data

```
csv = spark.read.csv('wasb:///data/flights.csv', inferSchema=True, header=True)
csv.show()
```

```
+----------+---------+-------+---------------+-------------+--------+--------+
|DayofMonth|DayOfWeek|Carrier|OriginAirportID|DestAirportID|DepDelay|ArrDelay|
+----------+---------+-------+---------------+-------------+--------+--------+
|        19|        5|     DL|          11433|        13303|      -3|       1|
|        19|        5|     DL|          14869|        12478|       0|      -8|
|        19|        5|     DL|          14057|        14869|      -4|     -15|
```

Step 3: Prepare the Data

```
data = csv.select("DayofMonth", "DayOfWeek", "OriginAirportID", "DestAirportID", "DepDelay", "ArrDelay")
data.show()
```

```
+----------+---------+---------------+-------------+--------+--------+
|DayofMonth|DayOfWeek|OriginAirportID|DestAirportID|DepDelay|ArrDelay|
+----------+---------+---------------+-------------+--------+--------+
|        19|        5|          11433|        13303|      -3|       1|
```

Example: Regression II

Step 4: Split the Data

```
splits = data.randomSplit([0.7, 0.3])
train = splits[0]
test = splits[1]
train_rows = train.count()
test_rows = test.count()
print "Training Rows:", train_rows, " Testing Rows:", test_rows
```

```
Training Rows: 1891270  Testing Rows: 810948
```

Step 5: Prepare the Training Data

```
assembler = VectorAssembler(inputCols = ["DayofMonth", "DayOfWeek", "OriginAirportID", "DestAirportID", "DepDelay"],
                            outputCol="features")
training = assembler.transform(train).select(col("features"), (col("ArrDelay").cast("Int").alias("label")))
training.show()
```

```
+--------------------+-----+
|            features|label|
+--------------------+-----+
|[1.0,1.0,10140.0,...|  -11|
|[1.0,1.0,10140.0    |   17|
```

Step 6: Train a Regression Model

```
lr = LinearRegression(labelCol="label",featuresCol="features", maxIter=10, regParam=0.3)
model = lr.fit(training)
print "Model trained!"
```

```
Model trained!
```

Example: Regression III

Step 7: Prepare the Testing Data

```
testing = assembler.transform(test).select(col("features"), (col("ArrDelay")).cast("Int").alias("trueLabel"))
testing.show()
```

```
+--------------------+---------+
|            features|trueLabel|
+--------------------+---------+
|[1.0,1.0,10140.0,...|       -5|
|[1.0,1.0,10140.0,...|      -25|
```

Step 8: Test the Model

```
prediction = model.transform(testing)
predicted = prediction.select("features", "prediction", "trueLabel")
predicted.show()
```

```
+--------------------+--------------------+---------+
|            features|          prediction|trueLabel|
+--------------------+--------------------+---------+
|[1.0,1.0,10140.0,...|-3.7579519642834542|       -5|
|[1.0,1.0,10140.0,...| -9.743936006015753|      -25|
```

What is Pipeline?

ML: PipelineModel

Train Data → **fit** → Pipeline → PipelineModel

PipelineModel

Binarizer → VectorAssembler → LinearRegressionModel

Test Data → **transform** → PipelineModel → Predictions

ML: Pipelines

A *Pipeline* is an estimator that contains stages representing a resusable workflow. Pipeline stages can be either estimators or transformers.

Pipeline

Transformer → Transformer → Estimator

Pipeline

Binarizer → VectorAssembler → LinearRegression

Pipeline

| Converting types | Creating features | Fitting and applying linear regression |

| Str 1 | Str 2 | ... | → | Val 1 | Val 2 | ... | → | Features | Label | → | Feat.. | Label | Pr... |

?

SQL Transformer   Vector Assembler   Lin Regression

SQL Transformer

Example: Pipeline

Step 1: Define the Pipeline

```
strIdx = StringIndexer(inputCol = "Carrier", outputCol = "CarrierIdx")
catVect = VectorAssembler(inputCols = ["CarrierIdx", "DayofMonth", "DayOfWeek", "OriginAirportID", "DestAirportID"],
                          outputCol="catFeatures")
catIdx = VectorIndexer(inputCol = catVect.getOutputCol(), outputCol = "idxCatFeatures")
numVect = VectorAssembler(inputCols = ["DepDelay"], outputCol="numFeatures")
minMax = MinMaxScaler(inputCol = numVect.getOutputCol(), outputCol="normFeatures")
featVect = VectorAssembler(inputCols=["idxCatFeatures", "normFeatures"], outputCol="features")
dt = DecisionTreeClassifier(labelCol="label", featuresCol="features")
pipeline = Pipeline(stages=[strIdx, catVect, catIdx, numVect, minMax, featVect, dt])
```

```
piplineModel = pipeline.fit(train)
print "Pipeline complete!"
```
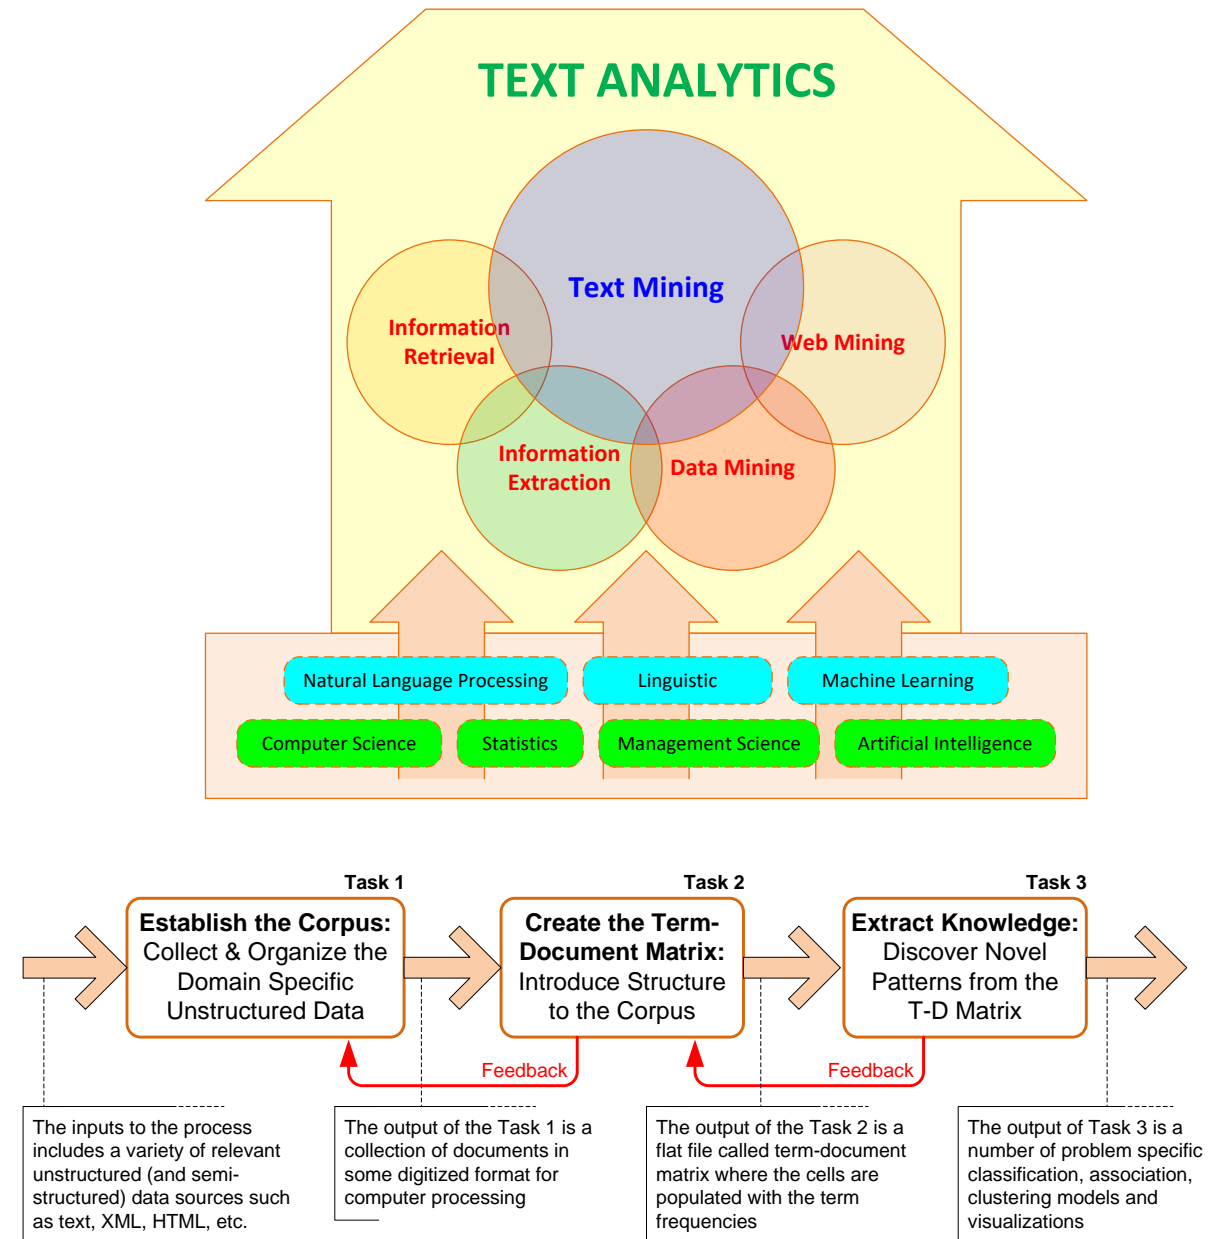
Step 2: Run the Pipeline as an Estimator

```
prediction = piplineModel.transform(test)
predicted = prediction.select("features", "prediction", "trueLabel")
predicted.show(100, truncate=False)
```

Step 3: Test the Pipeline Model

What is Text Analysis?

- 85-90 percent of all corporate data is in some kind of unstructured form (e.g., text)

- Unstructured corporate data is doubling in size every 18 months.

- Tapping into these information sources is not an option, but a need to stay competitive.

- Answer: text mining
  - A semi-automated process of extracting knowledge from unstructured data sources.
  - a.k.a. text data mining or knowledge discovery in textual databases.

TEXT ANALYTICS

Text Mining

Information Retrieval

Web Mining

Information Extraction

Data Mining

Natural Language Processing | Linguistic | Machine Learning

Computer Science | Statistics | Management Science | Artificial Intelligence

Task 1
Establish the Corpus: Collect & Organize the Domain Specific Unstructured Data

Task 2
Create the Term-Document Matrix: Introduce Structure to the Corpus

Task 3
Extract Knowledge: Discover Novel Patterns from the T-D Matrix

Feedback | Feedback

The inputs to the process includes a variety of relevant unstructured (and semi-structured) data sources such as text, XML, HTML, etc.

The output of the Task 1 is a collection of documents in some digitized format for computer processing

The output of the Task 2 is a flat file called term-document matrix where the cells are populated with the term frequencies

The output of Task 3 is a number of problem specific classification, association, clustering models and visualizations

Example: Text Analysis

Define the Pipeline

The pipeline for the model consist of the following stages:

- A Tokenizer to split the tweets into individual words.
- A StopWordsRemover to remove common words such as "a" or "the" that have little predictive value.
- A HashingTF class to generate numeric vectors from the text values.
- A LogisticRegression algorithm to train a binary classification model.

```
tokenizer = Tokenizer(inputCol="SentimentText", outputCol="SentimentWords")
swr = StopWordsRemover(inputCol=tokenizer.getOutputCol(), outputCol="MeaningfulWords")
hashTF = HashingTF(inputCol=swr.getOutputCol(), outputCol="features")
lr = LogisticRegression(labelCol="label", featuresCol="features", maxIter=10, regParam=0.01)
pipeline = Pipeline(stages=[tokenizer, swr, hashTF, lr])
```

```
piplineModel = pipeline.fit(train)
print "Pipeline complete!"
```

Step 2: Run the Pipeline as an Estimator

Step 3: Test the Pipeline Model

```
prediction = piplineModel.transform(test)
predicted = prediction.select("SentimentText", "prediction", "trueLabel")
predicted.show(100, truncate = False)
```

What is Clustering?

- Used for automatic **identification of natural groupings** of things.

- Part of the machine-learning family.

- Employ unsupervised learning.

- Learns the clusters of things from past data, then assigns new instances.

- There is not an output variable.

- Also known as **segmentation**.

Clustering results may be used to

- Identify natural **groupings of customers**.
- Identify rules for **assigning new cases to classes** for targeting/diagnostic purposes.
- Provide characterization, definition, labeling of populations.
- **Decrease the size and complexity** of problems for other data mining methods.
- **Identify outliers** in a specific domain (e.g., rare-event detection)

Clustering

Example: Clustering

Step 1: Create the K-Means Model

```python
assembler = VectorAssembler(inputCols = ["Age", "MaritalStatus", "IncomeRange", "Gender", "TotalChildren",
                            "ChildrenAtHome", "Education", "Occupation", "HomeOwner", "Cars"], outputCol="features")
train = assembler.transform(customers)

kmeans = KMeans(featuresCol=assembler.getOutputCol(), predictionCol="cluster", k=5, seed=0)
model = kmeans.fit(train)
print "Model Created!"
```

```python
centers = model.clusterCenters()
print("Cluster Centers: ")
for center in centers:
    print(center)
```

Step 2: Get the Cluster Centers

```python
prediction = model.transform(train)
prediction.groupBy("cluster").count().orderBy("cluster").show()
```

```python
prediction.select("CustomerName", "cluster").show(50)
```

Step 3: Predict Clusters

What is Recommendation?

Product recommender system

- Relevant product recommendations brings to a retailer the following value:
  - A client discovers new interesting products
    - → good customer experience
    - → more purchases
  - A client finds products he/she already know and like faster
    - → good customer experience
    - → more purchases

- The Netflix contest: Build a better recommender system from Netflix data

- Use the crowd's votes to complete the missing entries

	Aliens	Bug's Life	Cars	Dark Knight
Carmen	5	–	4	1
Joseph	5	4	–	–
Leonore	1	7	–	3
Esmerelda	2	8	1	–

Example: Recommendation I

Import the ALS class

In this exercise, you will use the Alternating Least Squa

```python
from pyspark.ml.recommendation import ALS
```

Load Source Data

The source data for the recommender is in two files - one containing numeric IDs for movies and users, a of the movies.

```python
ratings = spark.read.csv('wasb:///data/ratings.csv', inferSchema=True, header=True)
movies = spark.read.csv('wasb:///data/movies.csv', inferSchema=True, header=True)
ratings.join(movies, "movieId").show()
```

Prepare the Data

To prepare the data, split it into a training set and a test set.

```python
data = ratings.select("userId", "movieId", "rating")
splits = data.randomSplit([0.7, 0.3])
train = splits[0].withColumnRenamed("rating", "label")
test = splits[1].withColumnRenamed("rating", "trueLabel")
train_rows = train.count()
test_rows = test.count()
print "Training Rows:", train_rows, " Testing Rows:", test_rows
```

Example: Recommendation II

Build the Recommender

The ALS class is an estimator, so you can use its **fit** method to traing a model, or you can include it in a pipeline. Rath
label, the ALS algorithm requries a numeric user ID, item ID, and rating.

```
als = ALS(maxIter=5, regParam=0.01, userCol="userId", itemCol="movieId", ratingCol="label")
model = als.fit(train)
```

Test the Recommender

Now that you've trained the recommender, you can see how accurately it predicts known ratings in the test set.

```
prediction = model.transform(test)
prediction.join(movies, "movieId").select("userId", "title", "prediction", "trueLabel").show(100, truncate=False)
```

Questions